

Arjun Mehta

Bengaluru, India | arjun.mehta.dev@email.com | +91 98XXXXXX21 | arjunmehta.dev | github.com/arjunmehta

SUMMARY

Product-focused AI Engineer with 6+ years of experience building intelligent software products across Generative AI, machine learning, distributed systems, and modern web applications. Experienced in taking products from 0 → 1, designing scalable architectures, leading engineering teams, and converting ambiguous business problems into measurable technical solutions.

Currently focused on building AI systems that combine LLMs, retrieval-augmented generation, intelligent agents, and real-time analytics.

EDUCATION

Indian Institute of Technology, Hyderabad

B.Tech — Computer Science & Engineering, Relevant Coursework: Machine Learning · Artificial Intelligence · 2014 – 2018
Distributed Systems · Database Systems · Operating Systems · Computer Networks · Algorithms

EXPERIENCE

Lead AI Product Engineer

NOVA INTELLIGENCE LABS

Jan 2023 – Present

Bengaluru, India

Led the architecture and development of AI-powered products used by enterprise customers across multiple industries.

- Designed and deployed a RAG-based enterprise knowledge assistant processing more than 12 million internal documents.
- Reduced average information-retrieval time from 11 minutes to 42 seconds.
- Built an AI evaluation pipeline covering hallucination detection, retrieval accuracy, response quality, and latency.
- Designed multi-agent workflows for research, document analysis, and automated reporting.
- Reduced LLM infrastructure costs by 38% through caching, model routing, prompt optimization, and asynchronous processing.
- Led a team of 8 engineers across backend, frontend, ML, and DevOps.
- Worked directly with product leadership to convert customer requirements into technical roadmaps.
- Introduced automated CI/CD pipelines that reduced production deployment time from approximately 45 minutes to under 10 minutes.

Senior Software Engineer

FINTECHVERSE TECHNOLOGIES

Jul 2020 – Dec 2022

Pune, India

- Developed scalable backend services supporting more than 2.4 million monthly API requests.
- Designed event-driven microservices using Node.js and Python.
- Improved API response latency by 54% through database indexing, caching, and query optimization.
- Built real-time analytics dashboards using React and WebSocket-based data streams.
- Implemented automated monitoring and alerting for critical production services.
- Mentored junior engineers and conducted technical design reviews.
- Collaborated with security teams to improve authentication, authorization, and API security.

Software Engineer

PIXELFORGE DIGITAL

Aug 2018 – Jun 2020

Hyderabad, India

- Built customer-facing web applications using React, JavaScript, Node.js, and MongoDB.
- Developed reusable UI components used across five different products.
- Implemented REST APIs and authentication systems.
- Created automated testing workflows and improved overall release reliability.

- Worked closely with designers and product managers in Agile/Scrum environments.

PROJECTS

AURA — AI RESEARCH ASSISTANT | An AI-powered research platform capable of analyzing large document collections and generating source-grounded research reports.

- Technologies: Python, FastAPI, React, RAG, Vector Database, LLM APIs
- Impact: Reduced manual research time by approximately 70%.
- Supported document ingestion across PDF, DOCX, TXT, and web sources.
- Implemented citation-aware answer generation.
- Added evaluation framework for retrieval and response quality.

SENTINEL — REAL-TIME ANOMALY DETECTION | Machine-learning platform designed to detect abnormal patterns in high-volume transactional data.

- Technologies: Python, PyTorch, Kafka, PostgreSQL, Docker
- Result: Reduced average anomaly-detection response time from 8 minutes to under 30 seconds.

OPEN SOURCE AI TOOLKIT | Created a collection of reusable utilities for prompt testing, model evaluation, document chunking, embeddings, and RAG experimentation.

- Status: 1,800+ GitHub stars | 120+ forks

TECHNICAL SKILLS

Generative AI

LLMs

RAG

Prompt Engineering

AI Agents

NLP

Embeddings

Vector Search

Model Evaluation

Python

TypeScript

JavaScript

Java

React

Next.js

Node.js

FastAPI

REST APIs

Microservices

PyTorch

Scikit-learn

Pandas

NumPy

PostgreSQL

MongoDB

Redis

AWS

Docker

Kubernetes
GitHub Actions
CI/CD
Linux
OpenAI APIs
Hugging Face
LangChain
Pinecone
FAISS
MLflow

CERTIFICATIONS

AWS Certified Machine Learning – Specialty
NVIDIA Generative AI Certification
DeepLearning.AI — Generative AI with Large Language Models
Google Cloud — Professional Machine Learning Engineer

CAREER HIGHLIGHTS

6+ Years Engineering Experience 14+ AI/Software Products Delivered 2.4M+ Monthly API Requests Handled 38% Reduction in cloud infrastructure cost 47% Improvement in customer-support resolution time 3 Engineering Teams Led

SELECTED PRODUCT

Enterprise Knowledge Copilot Stack: Python · FastAPI · React · PostgreSQL · Redis · OpenAI API · Pinecone · AWS · Docker

LEADERSHIP & COMMUNITY

AI Community Mentor — 2023–Present Mentor early-career developers working on AI and software engineering projects.

Hackathon Judge — 2024 Evaluated student projects across AI innovation, technical implementation, scalability, and product usability.

Technical Speaker Delivered talks and workshops covering LLM applications, RAG architectures, and production AI engineering.

LANGUAGES

English — Professional Hindi — Native German — Basic

INTERESTS

Artificial Intelligence · Open Source · Chess · Product Design · Distributed Systems · Technology Communities

REFERENCES

Available upon request.